

Unstructured Data



卢浮部笃刻

关注



门萨智商测试题：以左下角的球为起点、右上角的球为终点，选择9个互相连接的球，并将球上的数字相加，能得出的最大值会是多少？



- A:457 B:51 C:其他

算法比答案更重要

这是从一道所谓门萨智商测试题修改而来的新题目。我想知道解这种题目的算法到底是什么样的。

除了暴力遍历，还有其他做法吗？

说点什么...

10

3

22

You are moving from the bottom-left ball to the top-right ball. Each round you can move up or move right, so you will pass 9 balls in total. Summing up these 9 numbers, what is the maximum value you can get?

Write an algorithm for it!

What's Unstructured Data?

- **Structured Data:** Data that has a standardized format for efficient access by software and humans alike.
- **Unstructured Data:** Information that either does not have a pre-defined data model or is not organized in a pre-defined manner.
- **Examples of Unstructured Data:** Text data (emails, news articles, books, reviews), images, videos, audio data etc.

Processing Text Data in Python

Processing Text Data in Python

#1: Word Frequency

Word Frequency

Given a document, you may want to find out the most frequently used words in the document; this is called word frequency analysis. It can be easily done with Python; but please remove the stop words (i.e., words such as “the”, “it”, “mine”, “but”) from the list!

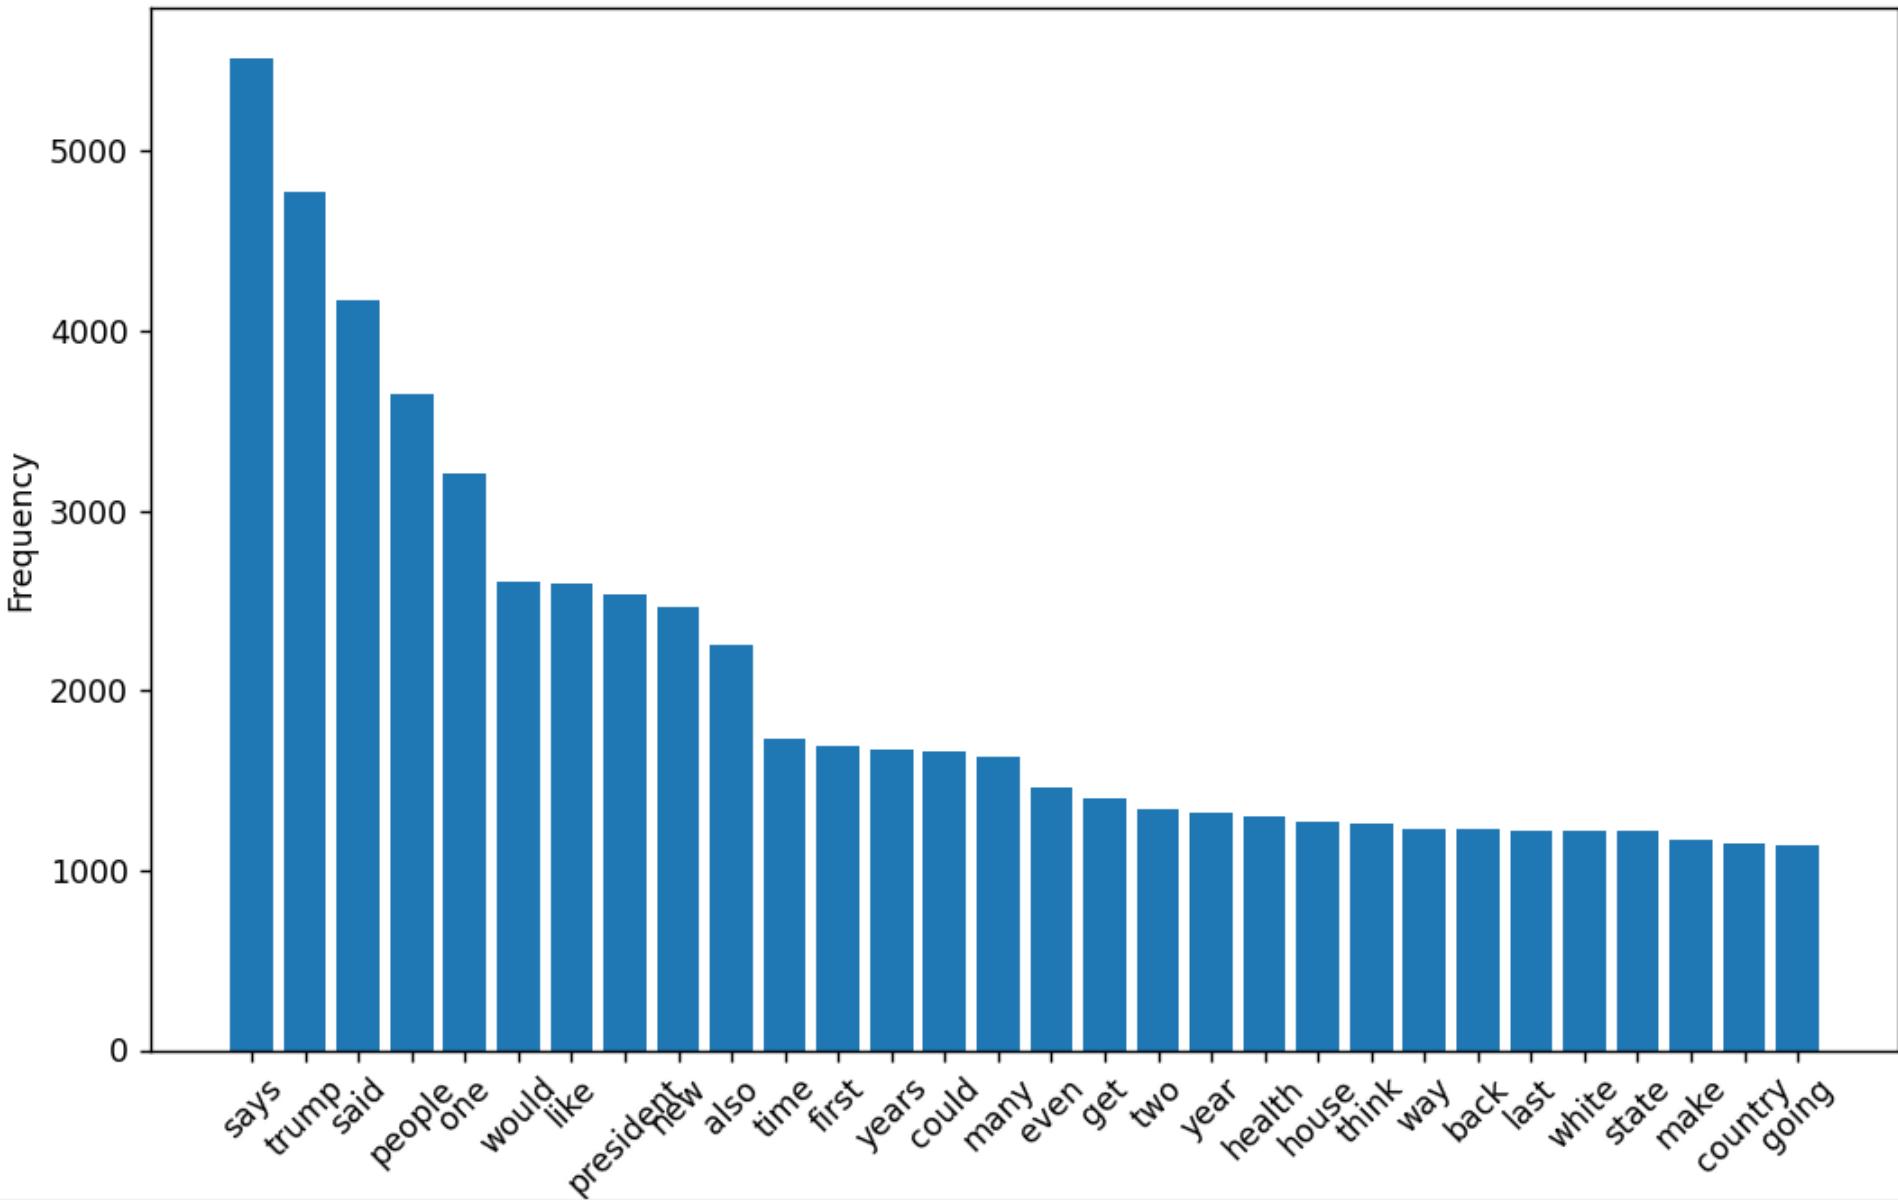
Word Frequency

Consider a text dataset which is available [here](#):

<https://ximarketing.github.io/data/LDAshort.txt>

Ask GPT to help you visualize the most frequent words used in the text document!

Top 30 Most Common Words



Processing Text Data in Python

#2: Topic Models

Topic Models

Topic modeling is an unsupervised learning algorithm. It takes a large number of documents (aka corpus) as input, and automatically generates a few topics from the text. Here, each topic is described by a distribution of words (typically we describe the most common words of the topic).

Topic Models

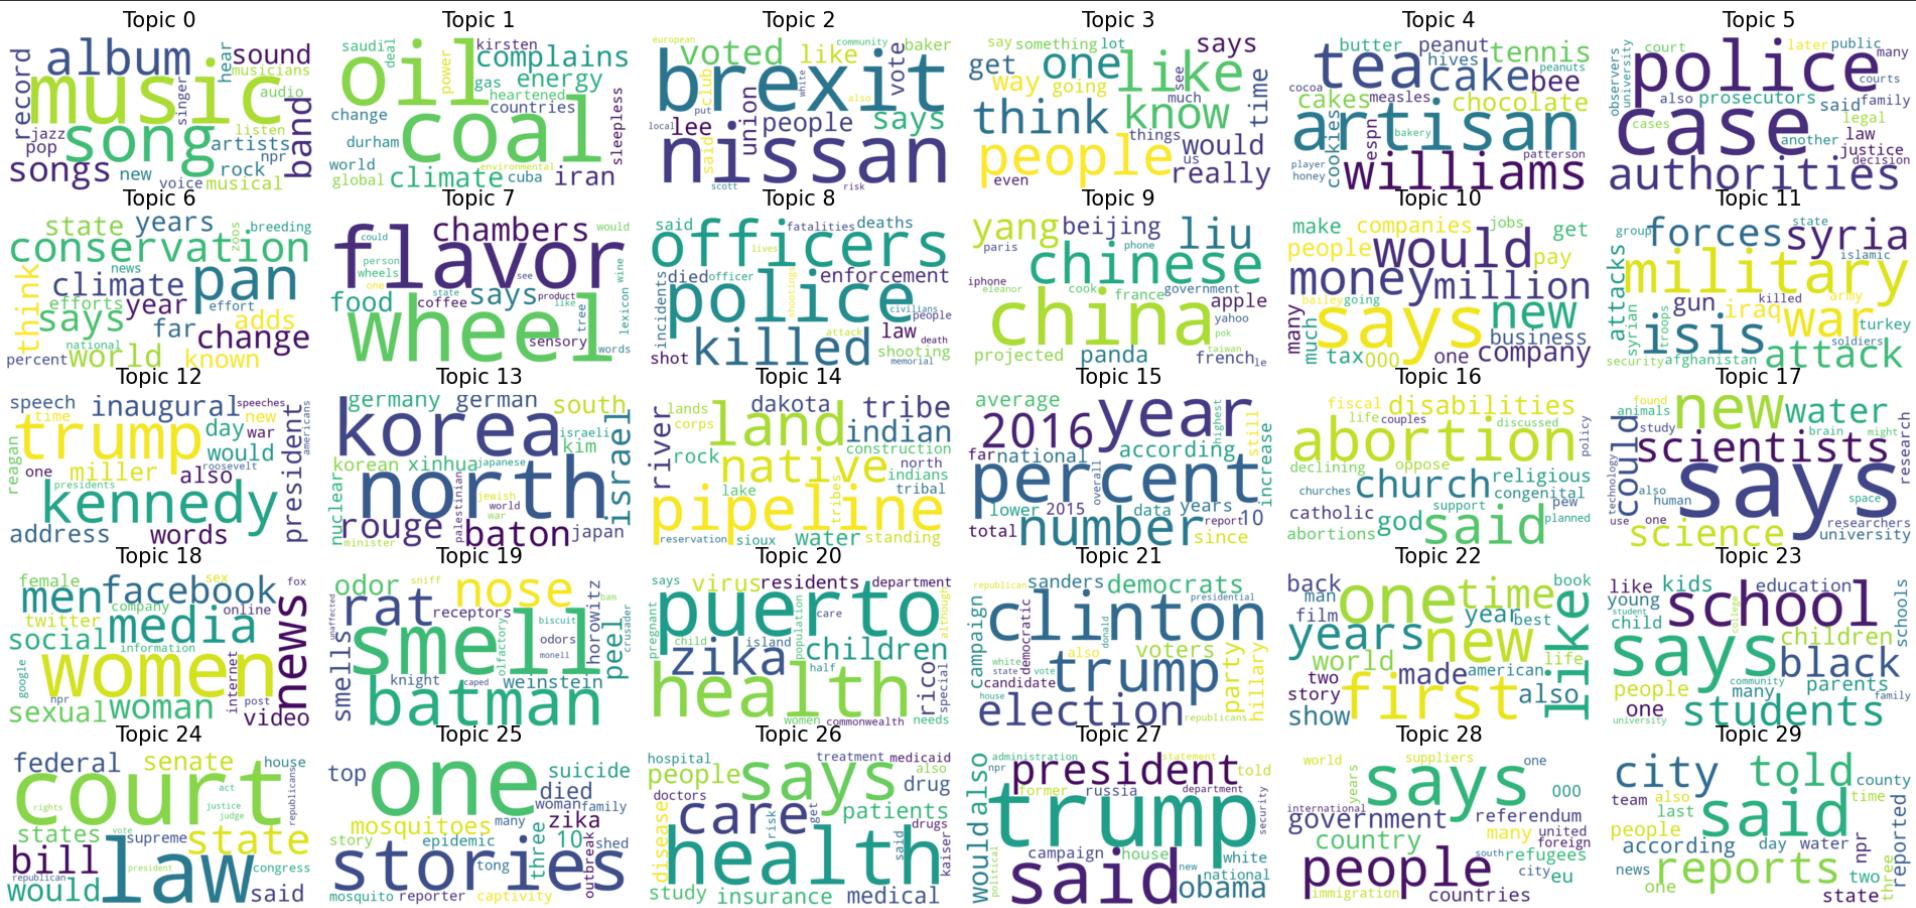
How does topic modeling work? The basic idea is, words of the same topic tend to appear simultaneously. For instance, in the topic that described “higher education”, some key words “university”, “library”, “examination”, “enrolment” often appear together.

You don't need to understand all these. Using topic models directly to get your results!

Topic Models

The most popular method for topic modeling is latent Dirichlet allocation (LDA). Let me show you an example based on LDA analysis of articles published by [Pew Research Center](#), an American think tank.

I did not want to read any of their articles, but I do want to know what topics they are discussing.



Interpreting the topics:

Topic 0: music, song, album, band, ...

Topic 1: oil, coal, energy, climate, ...

Topic 9: china, chinese, beijing, panda, ...

Topic 18: woman, man, female, sexual, ...

Topic 21: trump, clinton, election, sanders, ...

Topic 27: trump, obama, president, russia, ...

Let's try to play with LDA!

The original dataset is too large and I am unable to share it with you. Instead, we have a shorter version of the data with it stored at the following [location](#):

<https://ximarketing.github.io/data/LDAshort.txt>

Here, each line is an article and there are 1000+ articles.

Ask GPT to develop write code for you, and visualize your result (note: you need to install several modules).

```
1 import nltk, requests
2 from nltk.corpus import stopwords
3 from nltk.tokenize import word_tokenize
4 from gensim import corpora
5 from gensim.models import LdaModel
6 from wordcloud import WordCloud
7 import matplotlib.pyplot as plt
8
9 # Download NLTK data files (only the first time)
10 nltk.download('punkt')
11 nltk.download('stopwords')
12 nltk.download('punkt_tab')
13
14 file_url = 'https://ximarketing.github.io/data/LDAshort.txt'
15 response = requests.get(file_url)
16 if response.status_code == 200:
17     content = response.text
18 documents = content.split('\n')
19 stop_words = set(stopwords.words('english'))
20
21 def preprocess(text):
22     tokens = word_tokenize(text.lower())
23     tokens = [word for word in tokens if word.isalnum()]
24     tokens = [word for word in tokens if word not in stop_words]
25     return tokens
26
27 processed_documents = [preprocess(doc) for doc in documents]
28
29 dictionary = corpora.Dictionary(processed_documents)
30 corpus = [dictionary.doc2bow(doc) for doc in processed_documents]
31 num_topics = 30
32 lda_model = LdaModel(corpus, num_topics=num_topics, id2word=dictionary, passes=15)
33 for idx, topic in lda_model.print_topics(-1):
34     print(f'Topic: {idx} \nWords: {topic}\n')
```

Processing Text Data in Python

#3: Summarizing Text

A research team led by Professor Kevin Tsia, programme director of the Biomedical Engineering Programme under the Faculty of Engineering of the University of Hong Kong (HKU), has developed an artificial intelligence (AI)-driven imaging tool that enables speedy and precise diagnosis of cancer patients, greatly enhancing the effectiveness of their medical treatment.

In joint collaborations with HKU's Li Ka Shing Faculty of Medicine (HKUMed) and Queen Mary Hospital, the team headed by Professor Tsia, has successfully demonstrated the effective use of their latest generative AI method, the Cyto-Morphology Adversarial Distillation (CytoMAD), on lung cancer patients as well as drug tests. Combined with their proprietary microfluidic technology, CytoMAD allows fast and cost-effective 'label-free' imaging of human cells to help clinicians assess the patients' tumour at the precision of individual cells, and also determine whether the patients have the risk of metastasis.

CytoMAD uses AI to automatically correct cell imaging inconsistencies, enhance cell images, and extract previously undetectable information from cell images. Such all-round capability in CytoMAD ensures accurate and reliable downstream data analysis and diagnosis. CytoMAD's capabilities have the potential to revolutionize cell imaging for meaningful analysis of cell properties and related health and disease information.

▲ TOP

Here is a lengthy article about HKU research.
What does it talk about?

Text Summarization

Imagine that you want to understand the meaning of the text, but you do not really want to read the whole paragraphs to grasp its meaning. You just need a short summary of the text. This can be done with text summarization. Here is my output:

a team led by Professor Kevin Tsia has developed an artificial intelligence-driven imaging tool that enables speedy and precise diagnosis of cancer patients .

Abstractive Summarization Using Transformers

```
1 from transformers import pipeline
2 # Load the summarization pipeline
3 summarizer = pipeline("summarization")
4 # Sample text
5 text = """
6 Input your text here
7 """
8 # Generate a summary
9 summary = summarizer(text, max_length=40, min_length=15,
10 do_sample=False)
11 print(summary[0]['summary_text'])
```

Processing Text Data in Python

#4: Using APIs

This will be covered later in the class.

Processing Images in Python

#1: Display an Image

Image I/O and display

```
1 from PIL import Image, ImageFont, ImageDraw
2 from io import BytesIO
3 import requests
4 import matplotlib.pyplot as plt
5 import PIL.ImageStat as stat
6 url = "https://ximarketing.github.io/data/mainbuilding.jpg"
7 response = requests.get(url)
8 image = Image.open(BytesIO(response.content))
9 print(image.width, image.height, image.mode, image.format)
10 image.show()
```

If you are reading image from a local folder, you can replace line 6 with your path:

```
1 image = Image.open("../mainbuilding.jpg")
```

Processing Images in Python

#2: Color image to grayscale

A grayscale image only has two colors, black and white, and each pixel is given by a value $0 \leq x_i \leq 255$. When x_i is close to 0, the pixel is blacker, and when x_i is closer to 255, the pixel is whiter. We can use a simple formula to convert color image:

$$x_i = 0.299 \cdot r_i + 0.587 \cdot g_i + 0.114 \cdot b_i.$$

Color image to gray scale

```
from PIL import Image
import requests
from io import BytesIO
url = "https://ximarketing.github.io/data/mainbuilding.jpg"
response = requests.get(url)
img = Image.open(BytesIO(response.content))
gray_image = img.convert('L')
gray_image.save('gray_image.jpg')
gray_image.show()
```

How about adding colors to grayscale images?

This is much harder because we lose a lot of information when creating grayscale images: Given x_i , there are many combinations (r_i, g_i, b_i) that satisfy

$$x_i = 0.299 \cdot r_i + 0.587 \cdot g_i + 0.114 \cdot b_i.$$

It's not so great...



Colorized Image

Processing Images in Python

#3: Adjusting Brightness

How about adjusting brightness of images?

When the color of a pixel is $(0, 0, 0)$, it is black.

When the color of a pixel is $(255, 255, 255)$, it is white.

Suppose that the color of a pixel is (r, g, b) . If you move the color toward $(0, 0, 0)$, it becomes darker. If you move the color toward $(255, 255, 255)$, it becomes brighter.

This makes your image brighter.

```
1 from PIL import Image, ImageEnhance
2 import requests
3 from io import BytesIO
4 # URL of the image
5 url = 'https://ximarketing.github.io/data/mainbuilding.jpg'
6 # Fetch the image from the URL
7 response = requests.get(url)
8 img = Image.open(BytesIO(response.content))
9 image = img.point(lambda x: (255+x)/2)
10 image.show()
```

If you want darker, replace line 9 with:

```
1 image = img.point(lambda x: x/2)
```

Exercise: Develop an APP with the following functions

- The user can upload an image from his or her folder.
- There is a grayscale button. When a user clicks on the button, it converts the original image into grayscale and displays the grayscale image on a display panel to the user.
- Similarly, there is a brighter button, which makes the image brighter, and a darker button, which makes the image darker, and display the new image.

Processing Images in Python

#4: Changing colors

Making the Image Red

```
1 from PIL import Image, ImageEnhance
2 import requests
3 from io import BytesIO
4 # URL of the image
5 url = 'https://ximarketing.github.io/data/mainbuilding.jpg'
6 # Fetch the image from the URL
7 response = requests.get(url)
8 image = Image.open(BytesIO(response.content))
9 for i in range(image.width):
10     for j in range(image.height):
11         pixel = (image.getpixel((i, j)))
12         r = int((pixel[0] + 255)/2)
13         g = pixel[1]
14         b = pixel[2]
15         image.putpixel((i, j), (r, g, b))
16 image.show()
```

Negating the Image

```
1 from PIL import Image, ImageEnhance
2 import requests
3 from io import BytesIO
4 # URL of the image
5 url = 'https://ximarketing.github.io/data/mainbuilding.jpg'
6 # Fetch the image from the URL
7 response = requests.get(url)
8 image = Image.open(BytesIO(response.content))
9 for i in range(image.width):
10     for j in range(image.height):
11         pixel = (image.getpixel((i, j)))
12         r = 255 - pixel[0]
13         g = 255 - pixel[1]
14         b = 255 - pixel[2]
15         image.putpixel((i, j), (r, g, b))
16 image.show()
```

Processing Images in Python

#5: Rotation

Rotating the Image

```
1 from PIL import Image, ImageEnhance  
2 import requests  
3 from io import BytesIO  
4 # URL of the image  
5 url = 'https://ximarketing.github.io/data/mainbuilding.jpg'  
6 # Fetch the image from the URL  
7 response = requests.get(url)  
8 image = Image.open(BytesIO(response.content))  
9 image_t = image.transpose(Image.FLIP_LEFT_RIGHT)  
10 image_t.show()
```

Rotating the Image, 45 degree

```
1 image_45 = image.rotate(45)  
2 image_45.show()
```

Processing Images in Python

#6: Adding Information

Adding Text to Image

```
1 from PIL import Image, ImageEnhance, ImageDraw, ImageFont
2 import requests
3 from io import BytesIO
4 # URL of the image
5 url = 'https://ximarketing.github.io/data/mainbuilding.jpg'
6 # Fetch the image from the URL
7 response = requests.get(url)
8 image = Image.open(BytesIO(response.content))
9 draw = ImageDraw.Draw(image)
10 font = ImageFont.truetype("arial.ttf", 22)
11 draw.text((435, 93), "teaches me Python", font = font)
12 image.show()
```

Adding a Thumbnail

```
1 from PIL import Image, ImageEnhance, ImageDraw, ImageFont
2 import requests
3 from io import BytesIO
4 url = "https://ximarketing.github.io/data/trump2.jpg"
5 response = requests.get(url)
6 image_thumbnail = Image.open(BytesIO(response.content))
7 image_thumbnail.thumbnail((120,120))
8 url = "https://ximarketing.github.io/data/mainbuilding.jpg"
9 response = requests.get(url)
10 image = Image.open(BytesIO(response.content))
11 image.paste(image_thumbnail, (15, 55))
12 image.show()
```

Blending two images

Next, we are going to α -blend two images. The two images must have the same size and mode, and blending works as follows:

$$output = (1 - \alpha) \cdot image_1 + \alpha \cdot image_2$$

Code for blending Trump and HKU:

```
1 from PIL import Image, ImageFont, ImageDraw
2 from io import BytesIO
3 import requests
4 url = "https://ximarketing.github.io/data/trump2.jpg"
5 response = requests.get(url)
6 image1 = Image.open(BytesIO(response.content))
7 url = "https://ximarketing.github.io/data/mainbuilding.jpg"
8 response = requests.get(url)
9 image2 = Image.open(BytesIO(response.content))
10 image2 = image2.resize((image1.width, image1.height), Image.BILINEAR)
11 image = Image.blend(image1, image2, alpha = 0.3)
12 draw = ImageDraw.Draw(image)
13 font = ImageFont.truetype("arial.ttf", 50)
14 draw.text((350, 93), "Make HKU Great Again!",
15           font = font, fill='red')
16 image.show()
```

Processing Images in Python

#7: Color Channels

Separating colors in an image

```
1 from PIL import Image, ImageFont, ImageDraw
2 from io import BytesIO
3 import matplotlib.pyplot as plt
4 import requests
5 url = "https://ximarketing.github.io/data/parrot.jpg"
6 response = requests.get(url)
7 image = Image.open(BytesIO(response.content))
8 ch_r, ch_g, ch_b = image.split()
9 # split the RGB image into 3 channels: R, G and B
10 plt.imshow(ch_r, cmap = plt.cm.Reds)
11 plt.show()
12 plt.imshow(ch_g, cmap = plt.cm.Greens)
13 plt.show()
14 plt.imshow(ch_b, cmap = plt.cm.Blues)
15 plt.show()
```

Changing the color of your parrot!

```
1 from PIL import Image, ImageFont, ImageDraw
2 from io import BytesIO
3 import matplotlib.pyplot as plt
4 import requests
5 url = "https://ximarketing.github.io/data/parrot.jpg"
6 response = requests.get(url)
7 image = Image.open(BytesIO(response.content))
8 ch_r, ch_g, ch_b = image.split()
9 # split the RGB image into 3 channels: R, G and B
10 image = Image.merge('RGB', (ch_b, ch_g, ch_r))
11 image.show()
```



Processing Images in Python

#8: Compare Images

Discover the differences between two images:



Discover the differences between two images:

```
1 from PIL.ImageChops import difference
2 from PIL import Image
3 from io import BytesIO
4 import requests
5 url = "https://ximarketing.github.io/data/????.jpeg"
6 response = requests.get(url)
7 image1 = Image.open(BytesIO(response.content))
8 url = "https://ximarketing.github.io/data/????.jpeg"
9 response = requests.get(url)
10 image2 = Image.open(BytesIO(response.content))
11 image = difference(image1, image2)
12 image.show()
```



Question: What are the applications of image comparison?

Calculating Image Difference is very useful!

- Detecting and removing similar images from your smartphone storage.
- Video compression: If video frames are similar, we can just save their differences, not the entire images.
- Motion analysis: detecting anomalies in CCTV recording and sports game videos.
- Analyzing the information content of videos in computer vision.

We can help you relax :)

图侦民警王奇:为研判嫌疑人身份 有时连看8小时监控

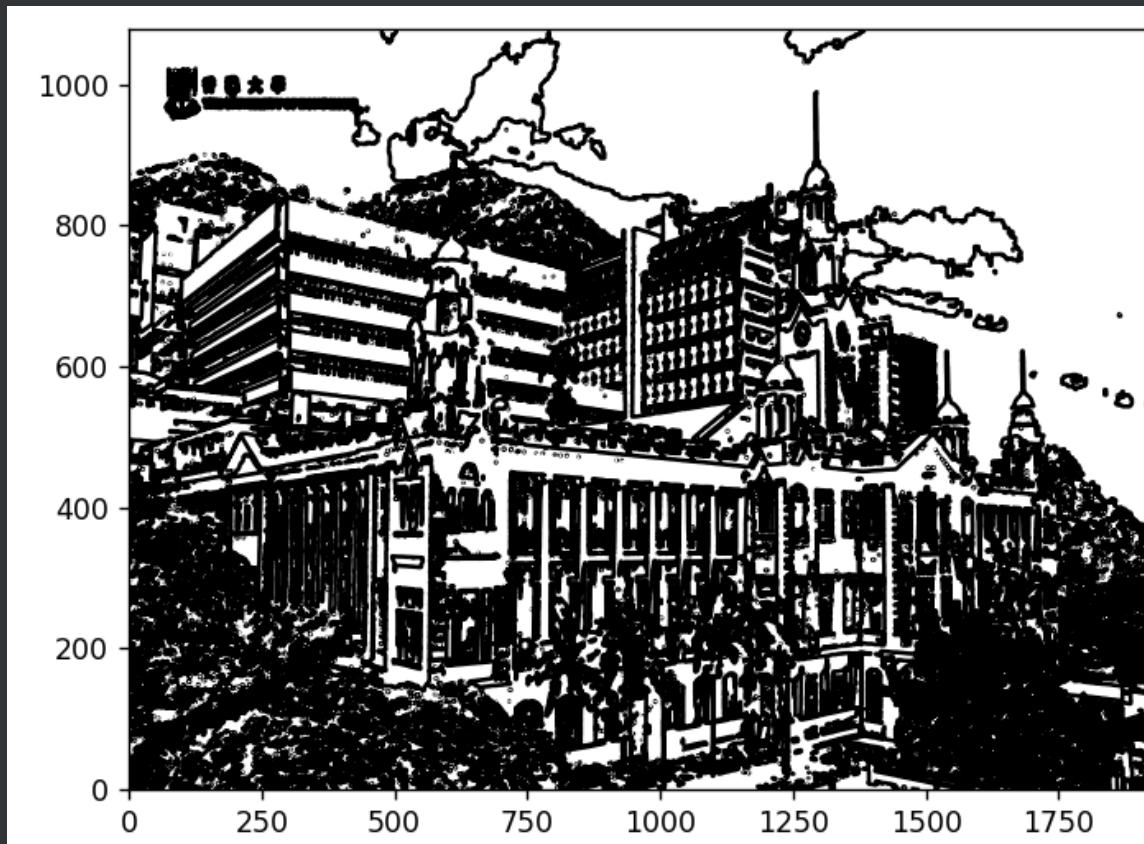
2022年02月10日 10:05 来源：新海南客户端

在街头遍布监控摄像头的时代，“图侦民警”逐渐被大众所熟悉，他们利用遍布街头巷尾的探头，通过观察、分析推演出犯罪分子的蛛丝马迹，为破案提供最直接的线索和可靠的依据。有人打比方说：“在图侦民警的手上，不会说话的摄像头成了无处不在的‘照妖镜’，科技的力量让犯罪无所遁形。”

Processing Images in Python

#9: Contour Line

How to turn your image into the following one?



Drawing Contour Line of an Image

A contour line for an image is a curve connecting all of the pixels where they have the grayscale value. We need to first convert an image into grayscale and then connect the pixels with the same grayscale intensity.

Drawing Contour Line of an Image

```
1 from PIL import Image
2 from io import BytesIO
3 import requests
4 import matplotlib.pylab as plt
5 import numpy as np
6 url = "https://ximarketing.github.io/data/trump2.jpg"
7 response = requests.get(url)
8 image = Image.open(BytesIO(response.content))
9 image = image.convert('L')
10 plt.contour(np.flipud(image), colors ='k',
11             levels = np.logspace(-15, 15, 100))
12 plt.show()
```

Processing Images in Python

#10: Convolution

What is convolution?

Convolution is an operation that operates on two images, one being an input image and the other one being a mask (also called the kernel) as a filter on the input image, producing an output image. It works by determining the value of a central pixel by adding the weighted values of all of its neighbors together to compute the new value of the pixel in the output image.

<https://www.youtube.com/embed/yb2tPt0QVPY?enablejsapi=1>

Why convolve an image?

Convolution applies a general-purpose filter effect on the input image. This is done in order to achieve various effects with appropriate kernels on an image, such as smoothing, sharpening, and embossing, and in operations such as edge detection.

Edge Detection

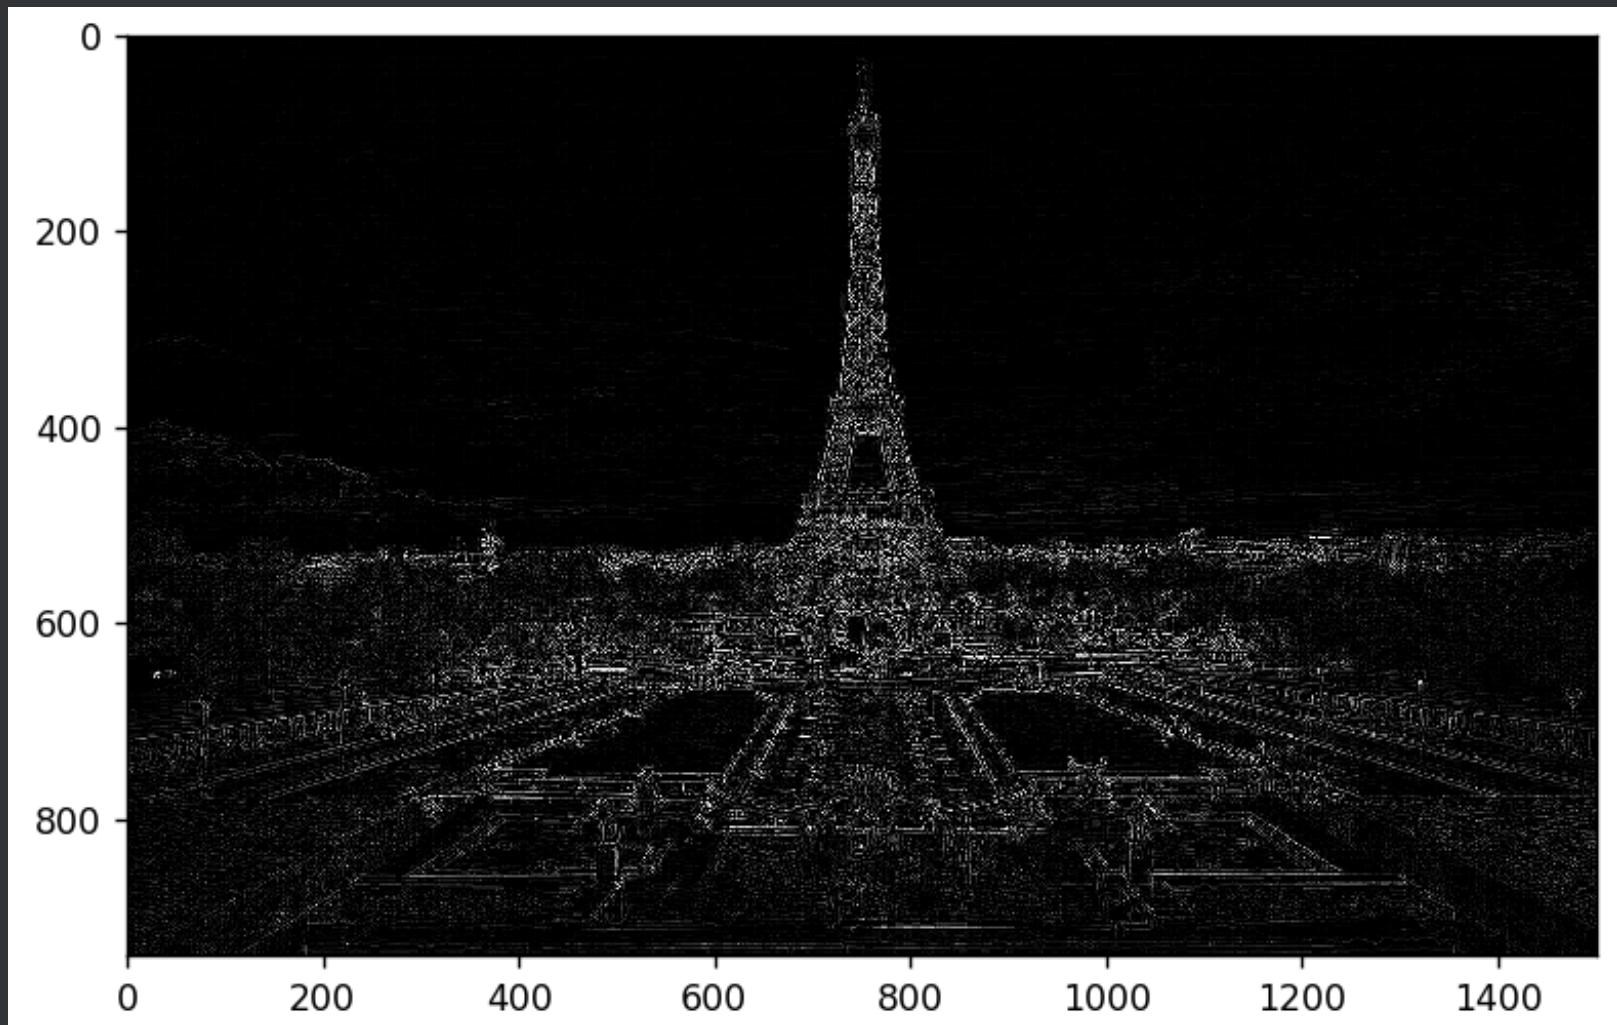
The following Laplace kernel is famous for edge detection. When applying the kernel to convolve images, you will get the edges of the original image.

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

Let's try the following code:

The convolve2d function allows you to convolve a 2D image.

```
1 from skimage import io
2 from skimage.io import imshow
3 from scipy import signal
4 from skimage.color import rgb2gray
5 import matplotlib.pyplot as plt
6 import numpy as np
7 url = 'https://ximarketing.github.io/data/effiel.jpg'
8 image = rgb2gray(io.imread(url)).astype(float)
9 edge_laplace_kernel = np.array([[0,1,0],[1,-4,1],[0,1,0]])
10 image_edges = np.clip(signal.convolve2d(image, edge_laplace_kernel), 0, 1)
11 imshow(image_edges)
12 plt.show()
```



Emboss an Image ("图像浮雕")

The following emboss kernel will be used. When applying the kernel to convolve images, you will get the embossed version of the original image.

$$\begin{bmatrix} -2 & -1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 2 \end{bmatrix}$$

Let's try the following code:

```
1 from skimage import io
2 from skimage.io import imshow
3 from scipy import signal
4 import matplotlib.pyplot as pylab
5 import numpy as np
6 url = 'https://ximarketing.github.io/data/effiel.jpg'
7 image = io.imread(url)/255 #
8 emboss_kernel = np.array([[-2,-1,0],[-1,1,1],[0,1,2]])
9 image_embossed = np.ones(image.shape)
10 for i in range(3):
11     image_embossed[...,i] = np.clip(
12         signal.convolve2d(image[...,i], emboss_kernel,
13                           mode='same', boundary="symm"), 0, 1)
14 imshow(image_embossed)
15 pylab.show()
```

More Kernels...

Kernel for Sharpening ("锐化")

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Box Blur Kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



Group Project (20%)

Four Options

1. Develop a game APP which demonstrates and solves a puzzle game. It features some algorithms that we learned in class (e.g., Hanoi tower).
2. Scraping some text data yourself (to be covered later) and perform some analysis of the text. In this case, an APP is not required.
3. Develop an image APP which allows users to upload images and perform some operations on the image.
4. Develop an APP which helps users scrape data.

Submission

1. Your Python code
2. Your **complete** chat history with AI
3. A user manual of the APP (for options 1,3 and 4) or a report of your findings (for option 2)
4. Supplementary data files, if needed

Deadline: 23:59 PM, April 19, 2025

Next Class

Each group demonstrates the APP you have developed for the previous assignment [data analysis and visualization] to the class, followed by a **peer evaluation**.

1~2 minutes for each group demonstration.